olist Data Engineering Lab

E-COMMERCE DATA PIPELINE & DATA INSIGHTS

Data Engineering Team:
- Bandekar Anjali Ranjit
- Irene Lua
- M Logaraj
- Yee Soon Tuck

olist Data Engineering Lab

Agenda

1. Project Management Approach
2. Logical Data Pipeline Architecture
3. Platform Data Pipeline Architecture (2 Variations)
4. Data Warehouse Design
5. ELT Pipeline
6. Data Quality Testing
7. Data Analysis & Insights
8. Pipeline Orchestration

olist Data Engineering Lab

Trello

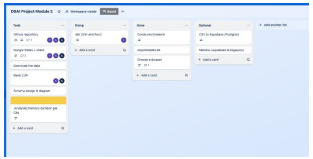

Project Management Strategy - Trello

- Trello was utilized for tracking project tasks and progress.
- Each team member has designated cards for their responsibilities.
- The board is organized into columns for 'To Do', 'In Progress', and 'Done'.
- Guide for daily meetings
- Able to give comments ensure clear communication and accountability

Benefits of Using Trello

- Visual organization enhances team collaboration and transparency.
- Easy to assign tasks and set deadlines within the platform.
- New ideas/links could be easily updated and read by team members
- Trello's mobile app allows team members to update tasks on the go.

olist Data Engineering Lab

Source Data

Data set: Brazilian E-Commerce Public Dataset by Olist

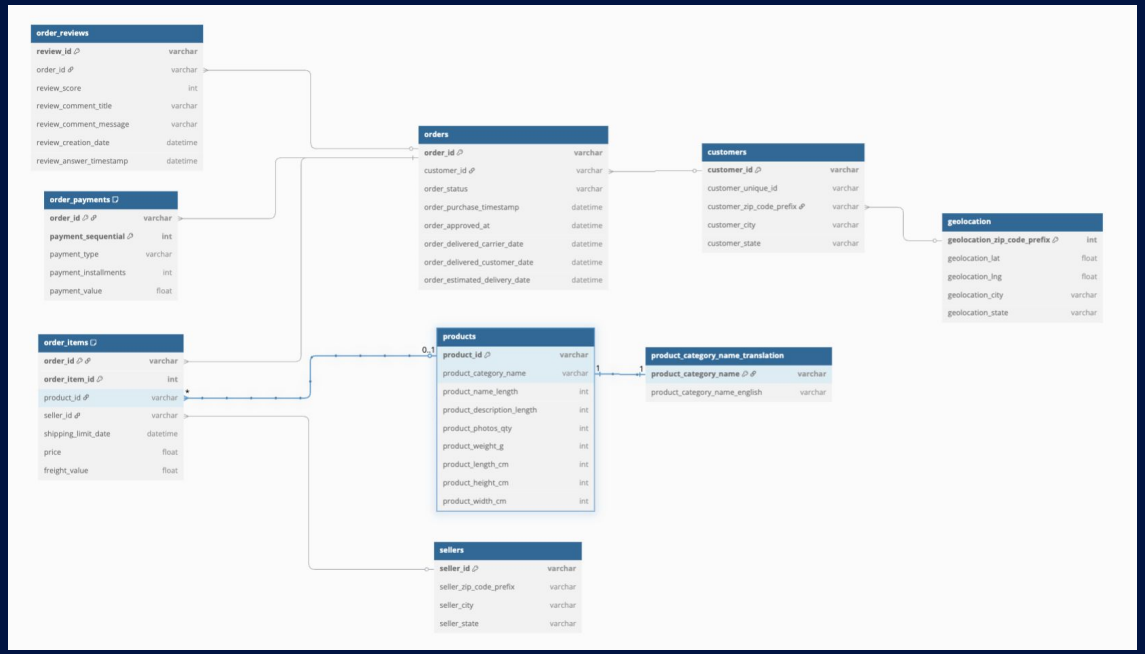

olist Data Engineering Lab

Logical Data Pipeline Architecture

Data Source

CSV Files

Data Storage

Data Buckets

Data Preparation

Data Warehouse

ETL

Analyses

Python

Pandas

Seaborn

olist Data Engineering Lab

Platform Data Pipeline Architecture v 1.0

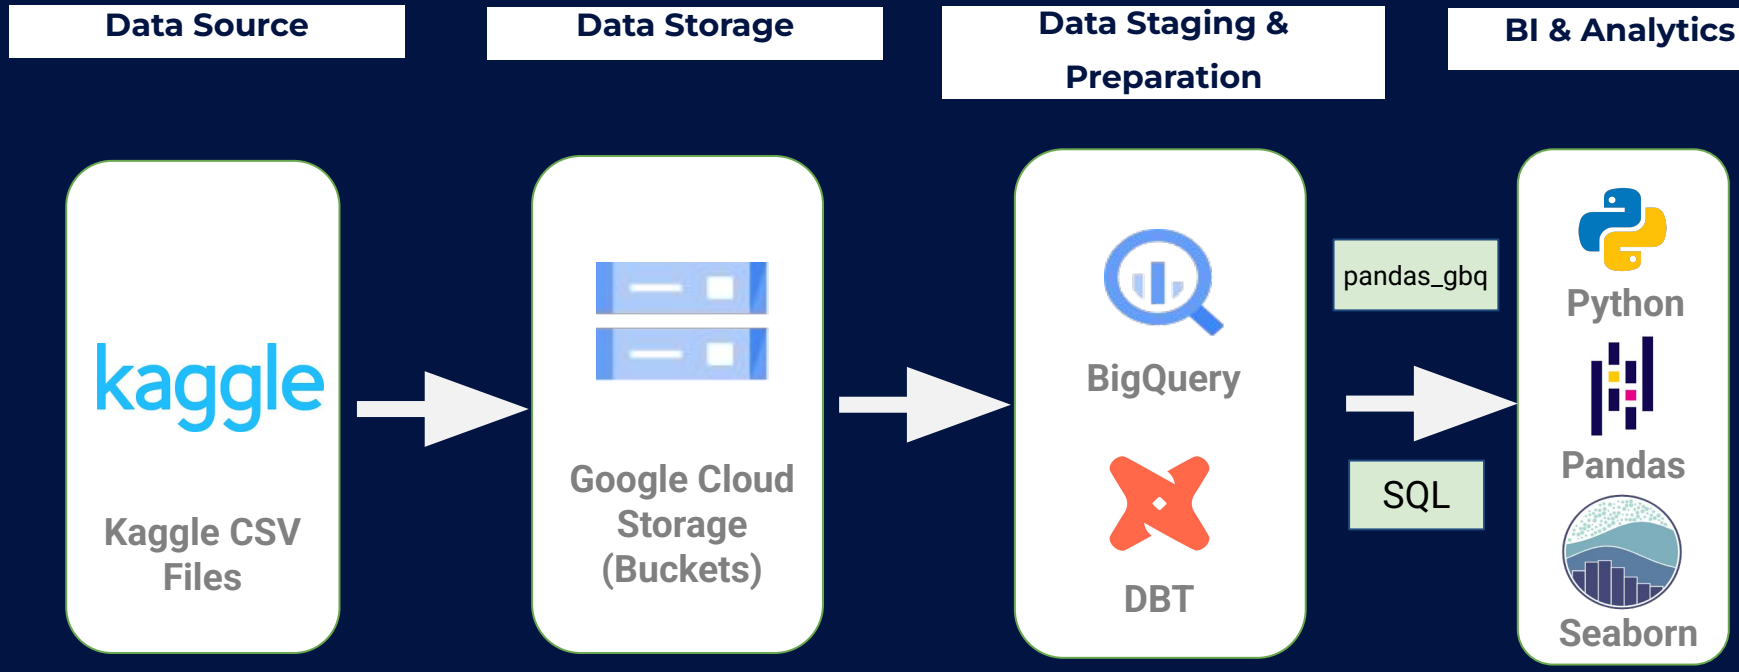

olist Data Engineering Lab

Platform Data Pipeline Architecture v 2.0

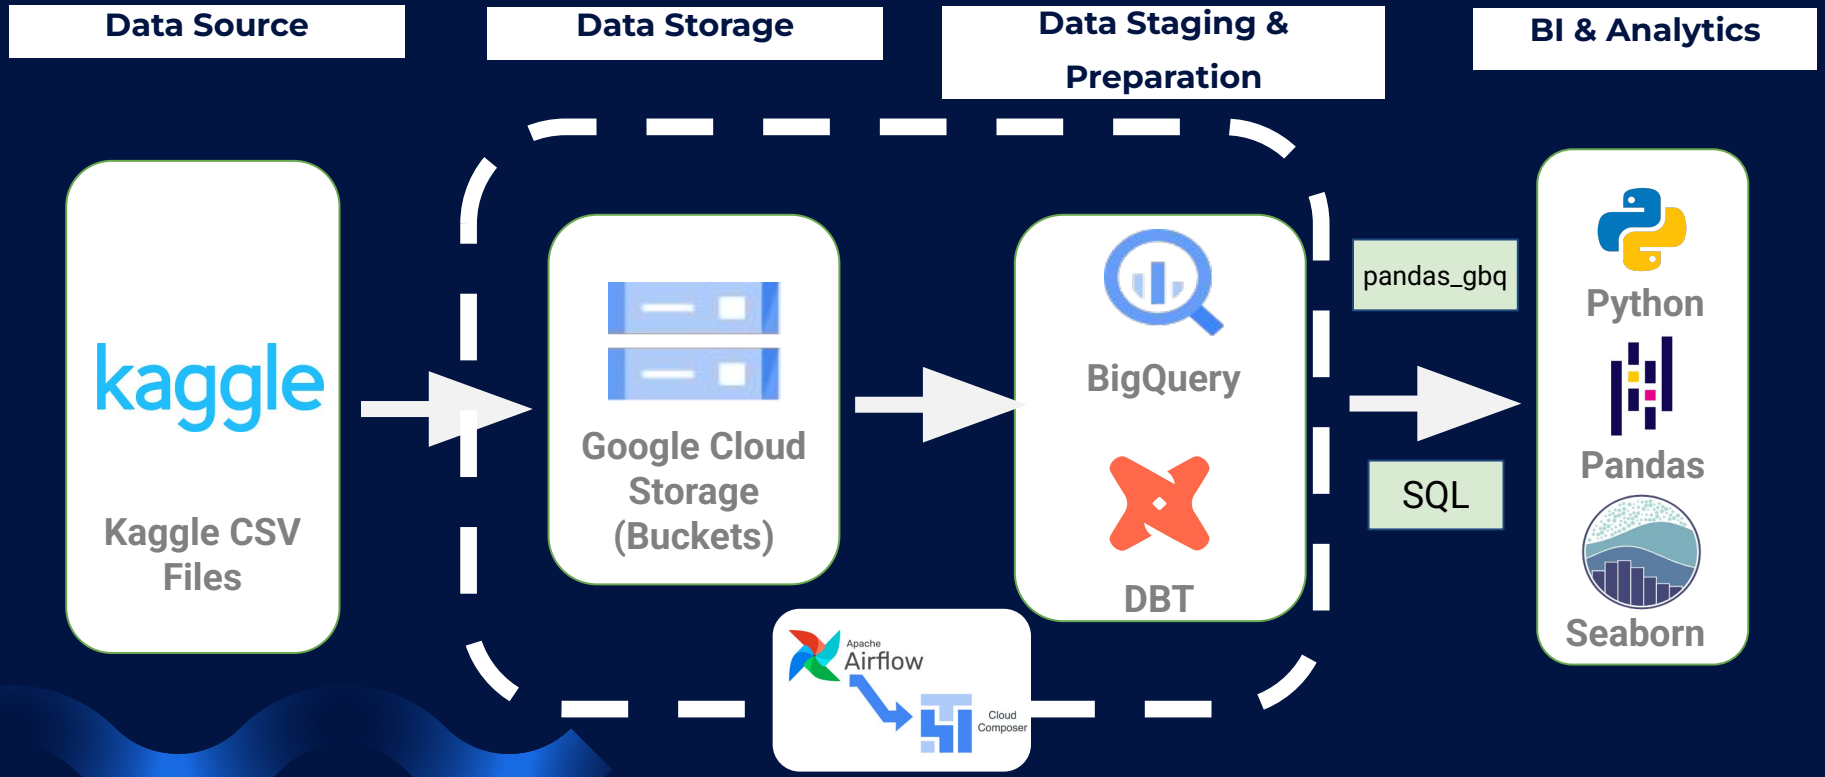

olist Data Engineering Lab

Data Preparation

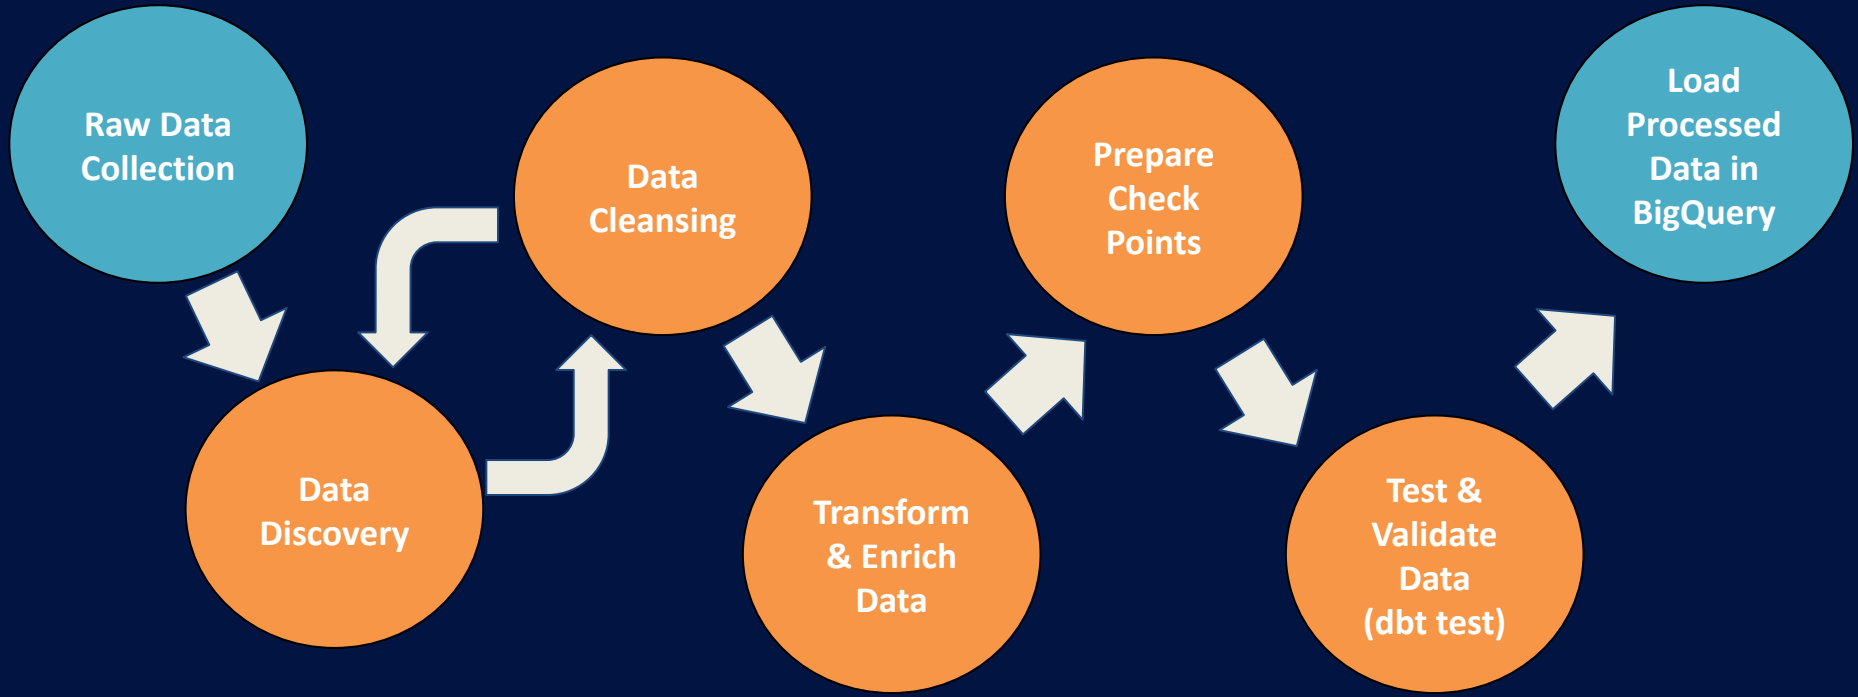

olist Data Engineering Lab

Test Case:

Does the delivered order value match our payment receipts?

	123 total_payment
1	16,008,872.119998764

A-Z order_status	123 count(order_status)	123 sum(price)
unavailable	643	[NULL]
canceled	179	[NULL]
created	5	[NULL]
delivered	3	134.97
invoiced	2	[NULL]
shipped	1	[NULL]
[NULL]	0	[NULL]

Unavailable?
Need to find out if link to deliveries.

123 total_price	123 total_freight	123 total_payment
13,591,643.7	2,251,909.54	15,846,280.17

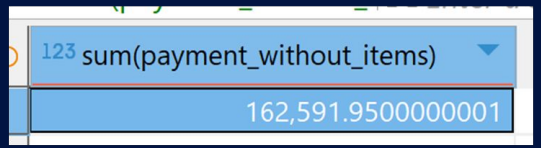

123 sum(payment_without_items)
162,591.9500000001

rows: 766 orders were paid without order items

```yaml
models:
  - name: check_order_item_payment
    description: 'Model description'
    columns:
      - name: order_status
        description: Only 'canceled' status is acceptable here.
        tests:
          - dbt_utils.expression_is_true:
              expression: "{{ order_status }} = 'canceled'"
```

Here is the content for your review first, I want to add nicer snapshots to illustrate the process.

olist Data Engineering Lab

Test Case:

Does the delivered order value match our payment receipts?

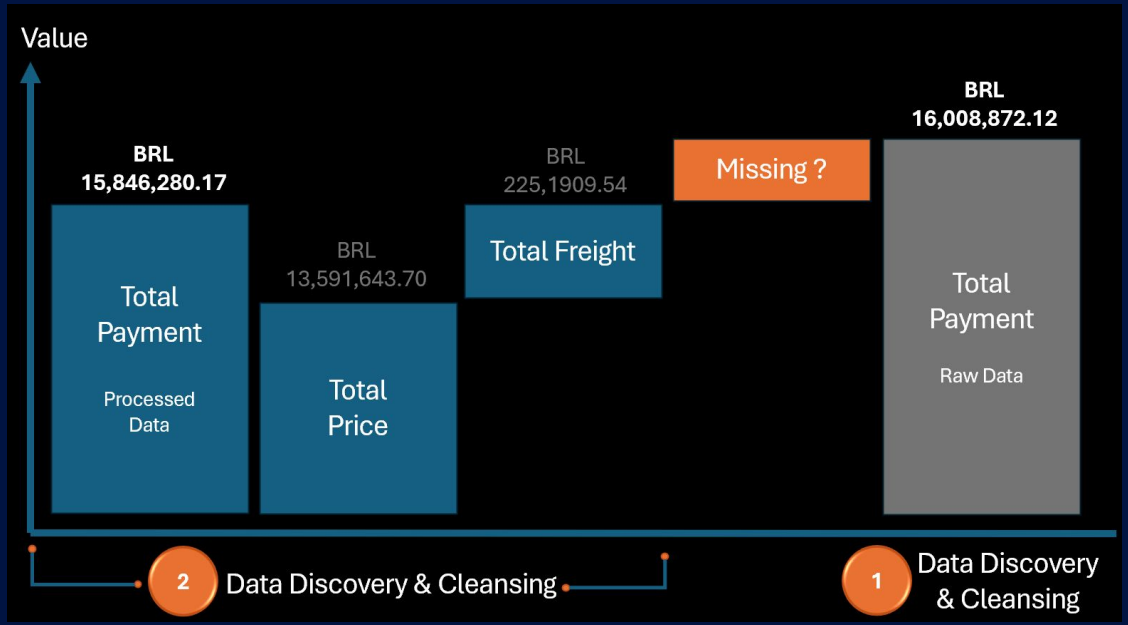

Test Case 1 working file link

New Query : Checking Missing Data

	order_status	order_status_count	price_sum
0	canceled	179	0.00
1	created	5	0.00
2	delivered	3	134.97
3	invoiced	2	0.00
4	shipped	1	0.00
5	unavailable	643	0.00

```yaml
models:
  - name: check_order_item_payment
    description: 'Model description'
    columns:
      - name: order_status
        description: Only 'canceled' status is acceptable here.
        tests:
          - dbt_utils.expression_is_true:
              expression: "{{ order_status }} = 'canceled'"
```

olist Data Engineering Lab

Data Warehouse Design (Star Schema)

Business Question: Providing Insights on Delivery Performance of Olist

1. Central Table — **fact_orders**
   a. Tracks every order placed, including key timestamps, delivery status, and total value.
      ➤ It's the heart of the schema where all delivery metrics are measured.
2. Customer Dimension — **dim_customers**
   a. Adds geographic context (city, state) to each order, enabling analysis of regional delivery trends and customer behavior.
3. Item Dimension — **dim_order_items**
   a. Breaks down orders into individual items with price, freight cost, and product category — crucial for understanding delivery cost drivers and product-based delays.
4. Payment Dimension — **dim_order_payments**
   a. Links payment methods and values to orders, allowing insights into how payment types affect delivery performance (e.g., delays with boleto payments).

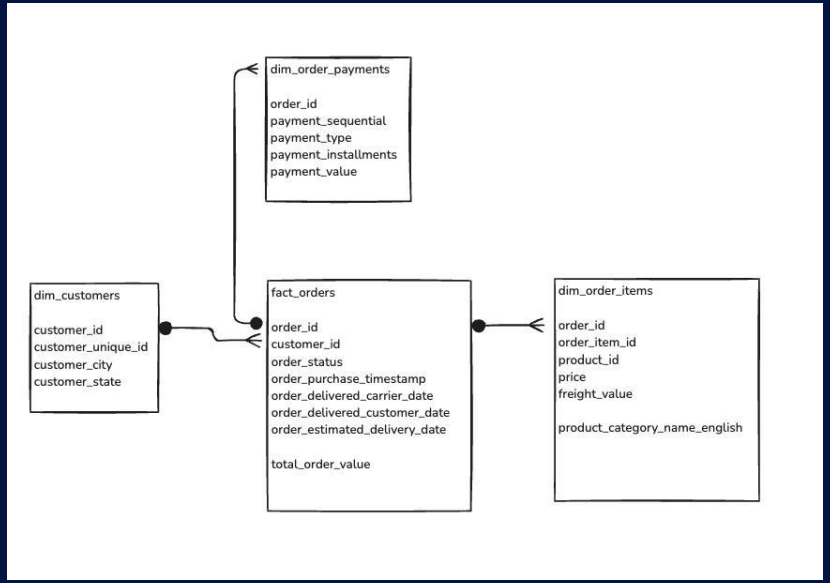

olist Data Engineering Lab

ELT Pipeline Features

1. **Transformation of Raw Data into a Star Schema Using dbt:**
   - Utilized dbt to structure the raw data into a star schema, enhancing query performance and simplifying data analysis.
   - Created SQL models to represent different dimensions and fact tables within the schema.

2. **Implementation of Data Cleaning and Validation Steps:**
   - Employed dbt's testing capabilities to ensure data quality by defining tests within the project, such as checking for uniqueness, non-null constraints, dbt_utilis functions

3. **Creation of Derived Columns:**
   - Developed additional metrics like **total_order_amount** within dbt models to enrich the dataset and provide deeper analytical insights.

4. **Exploratory Data Analysis Using Jupyter Notebooks:**
   - Utilized Jupyter Notebooks to perform exploratory data analysis on the transformed data, allowing for interactive examination and visualization of trends and patterns.

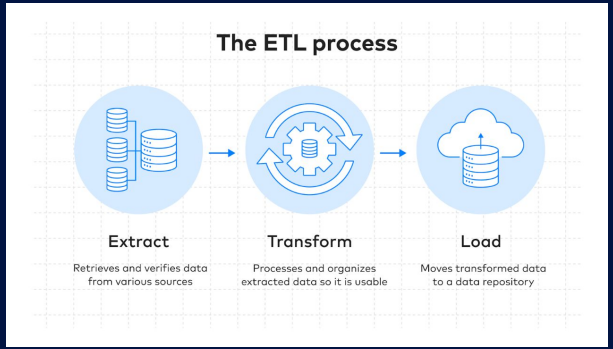

olist Data Engineering Lab

Data Analysis Insights

1) Northern States in Brazil had the highest average delivery dates

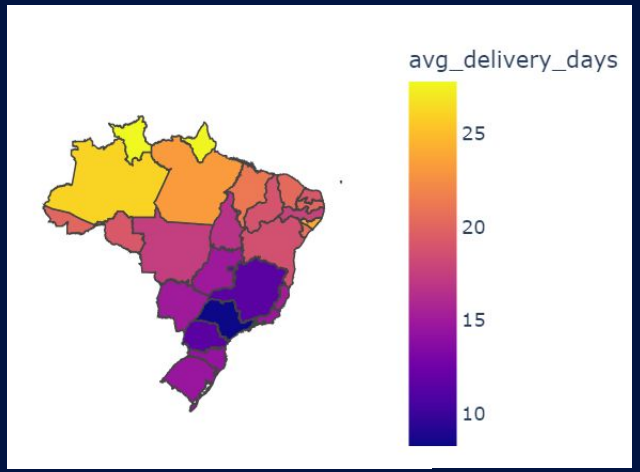

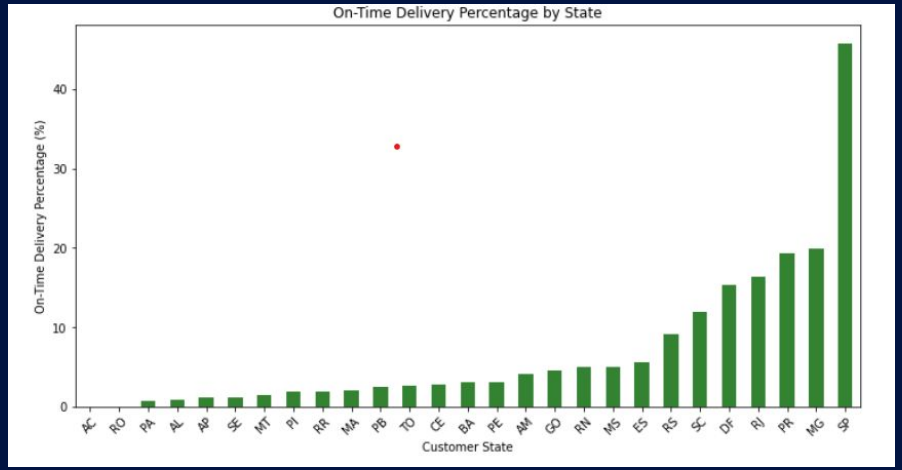

2) Bulkier Items show a stronger correlation to higher delivery times

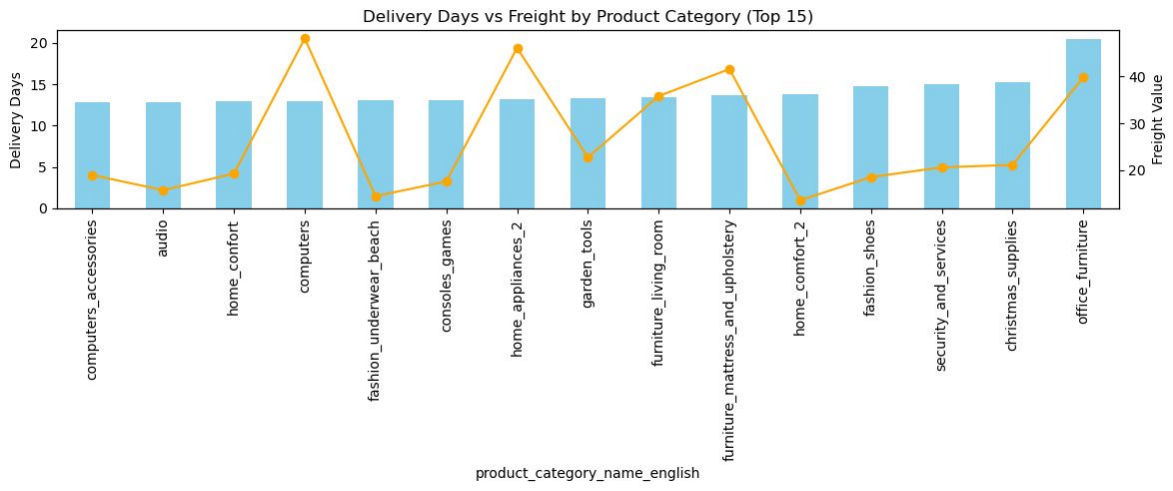

olist Data Engineering Lab

3) Delivery times peak in February, March, November, and December due to seasonal demand, holidays

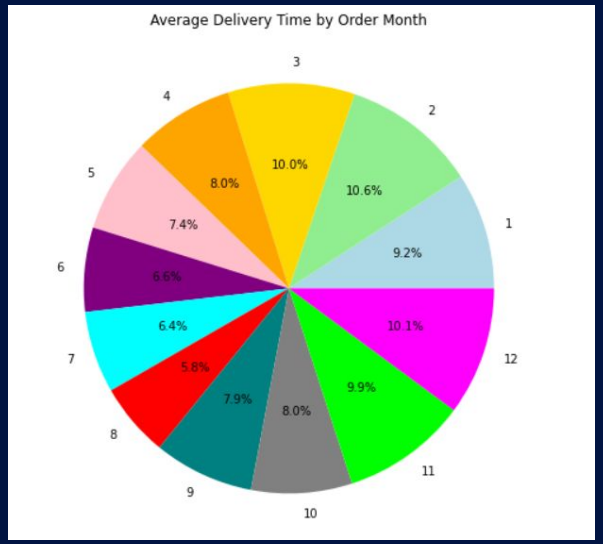

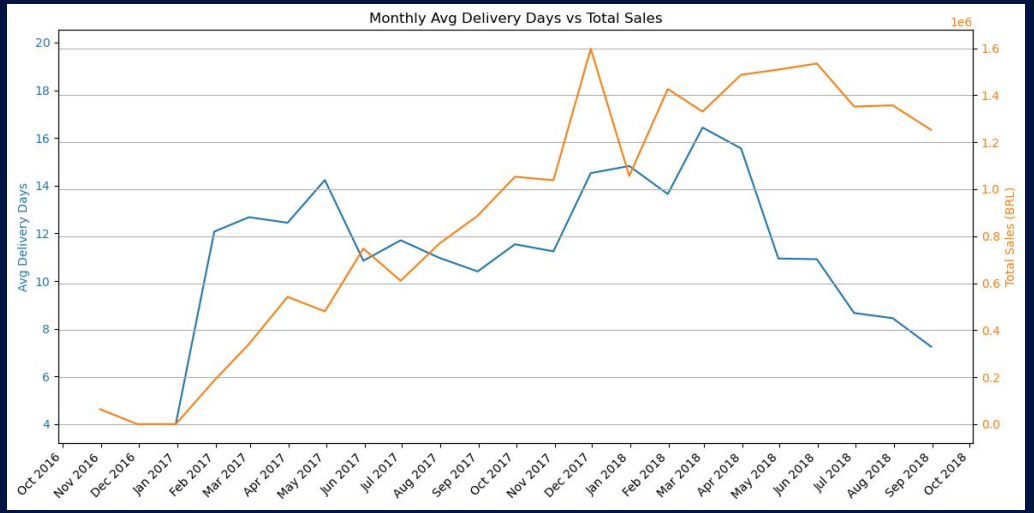

olist Data Engineering Lab

4) Boleto & Credit Card Transactions have higher delivery times

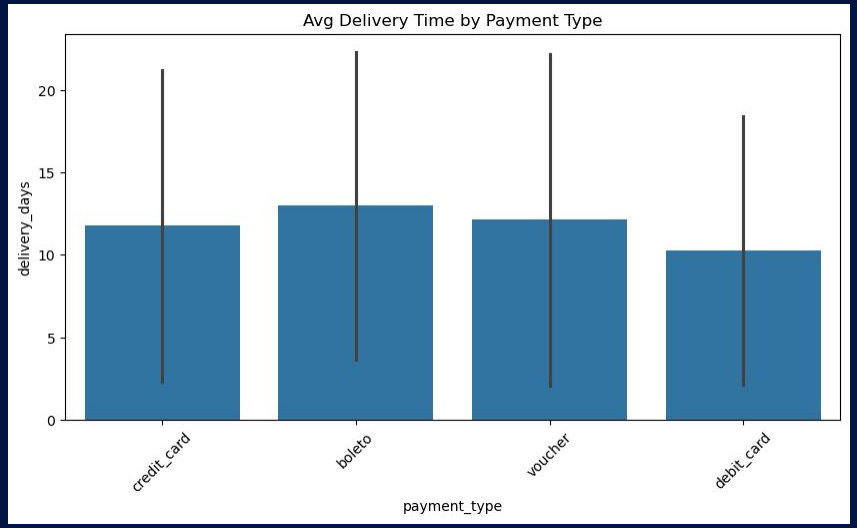

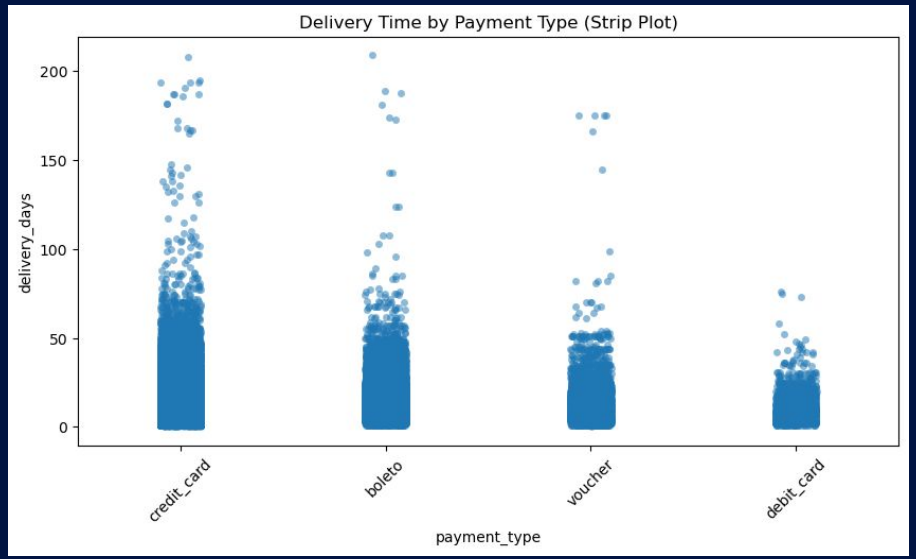

olist Data Engineer Lab

Pipeline Orchestration - Google Cloud Composer

Managed service, using Apache Airflow.
Note: this orchestration pipeline is connected to a different dataset in BigQuery (composer_test) for the purpose of presentation.
The other parts of this Project uses the “brazil_e_commerce” dataset.

olist Data Engineer Lab

Data Pipeline in Google Cloud Composer

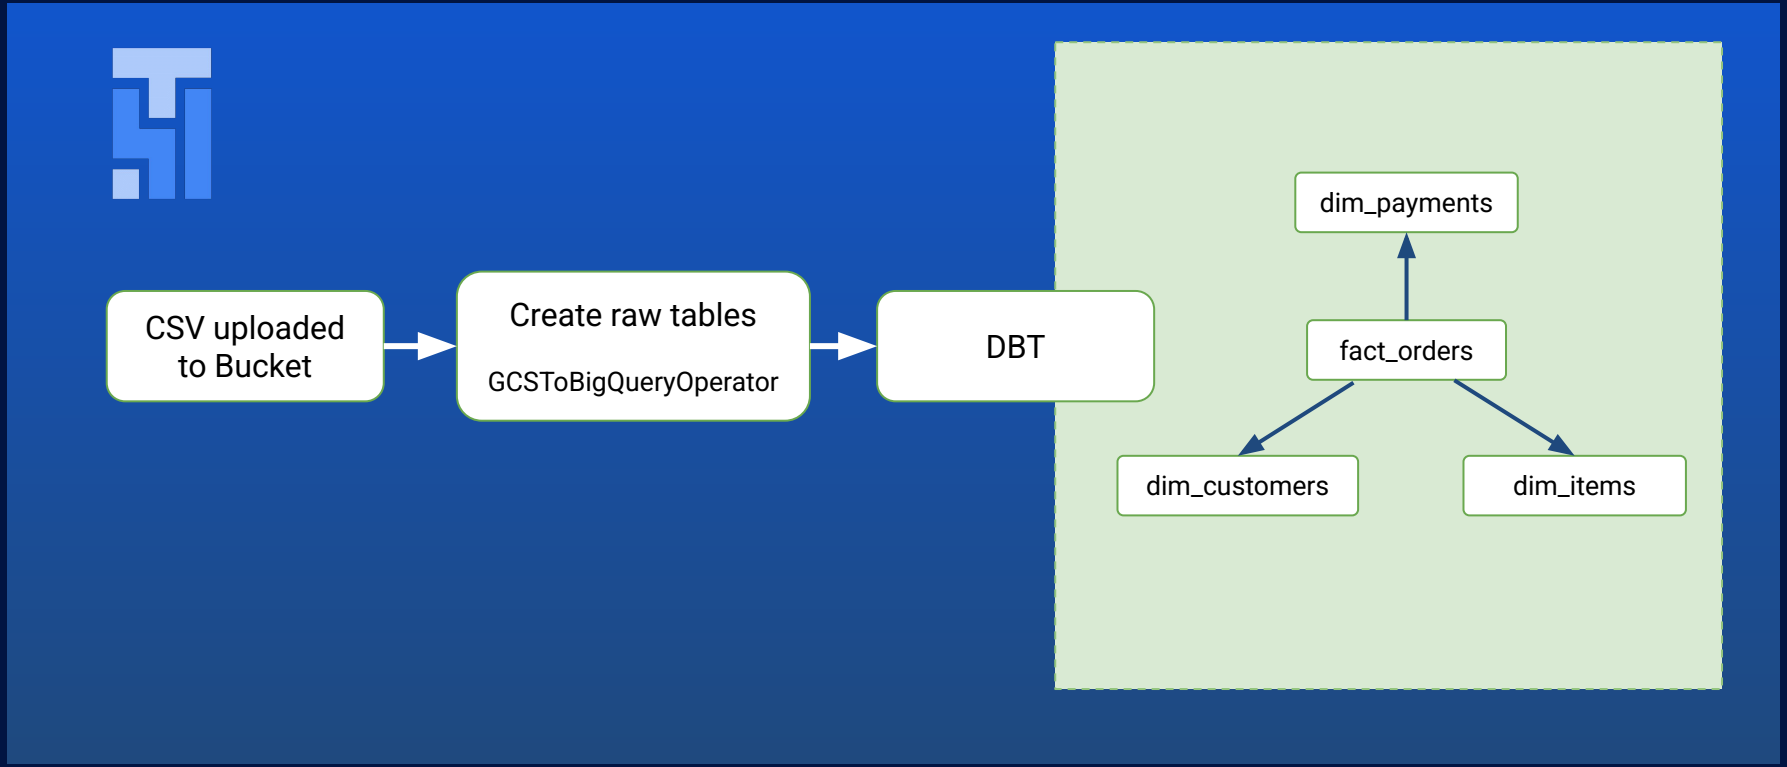

olist Data Engineer Lab

DAGs folder in Bucket

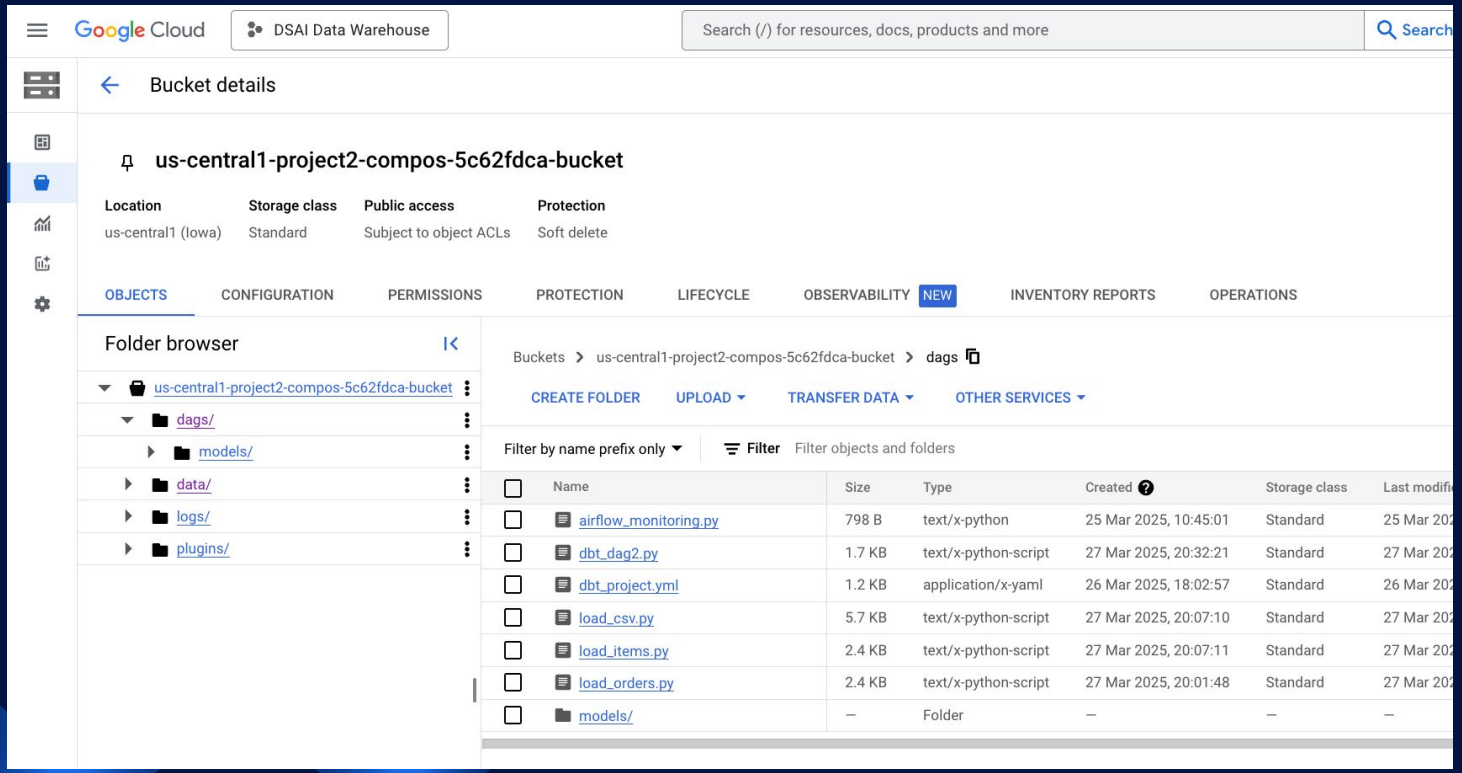

olist Data Engineer Lab

Airflow DAGs in Google Cloud Composer

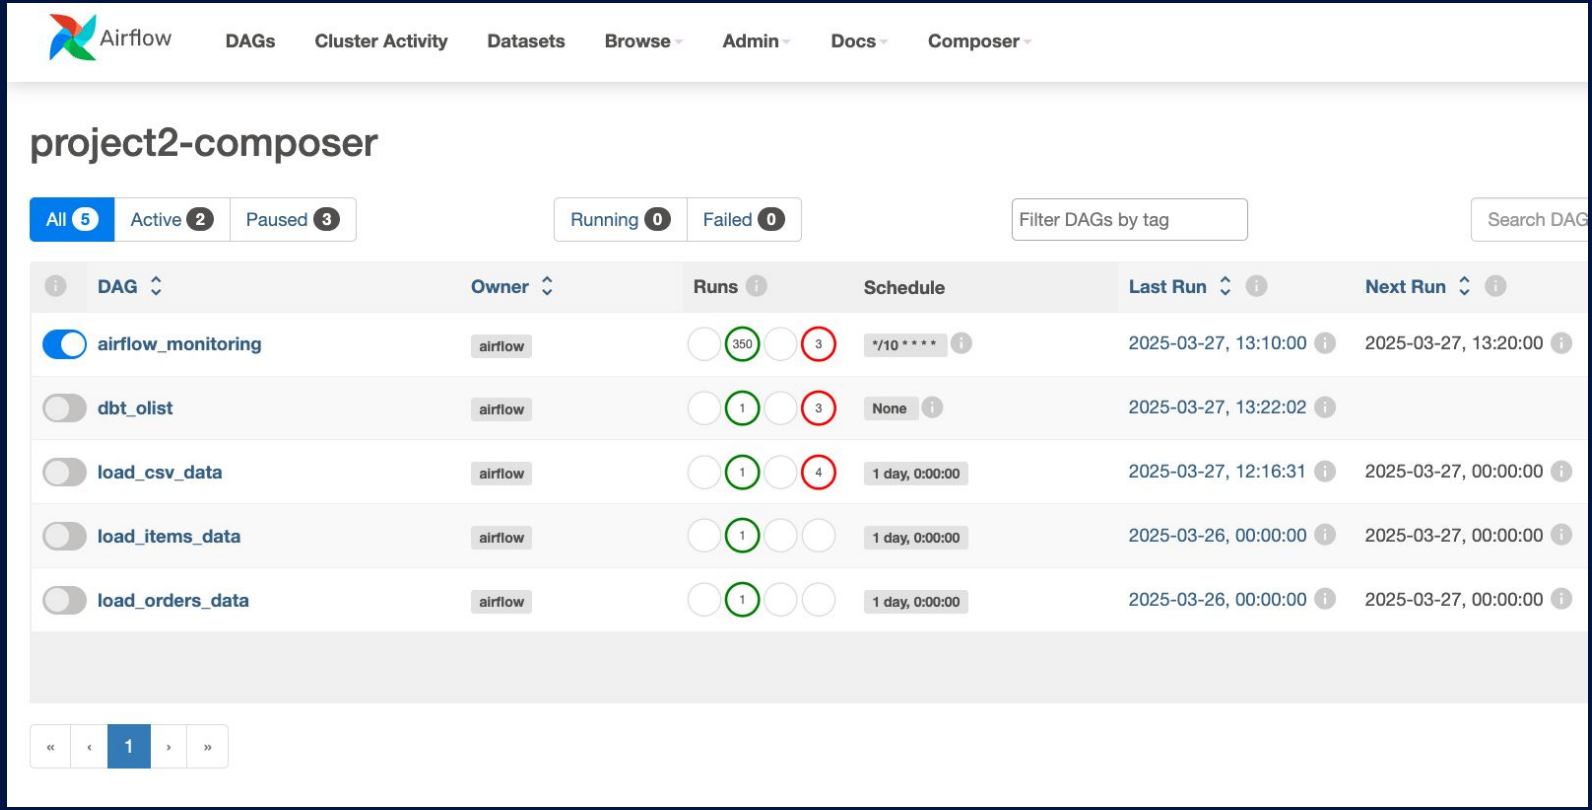

load_csv_data DAG run logs

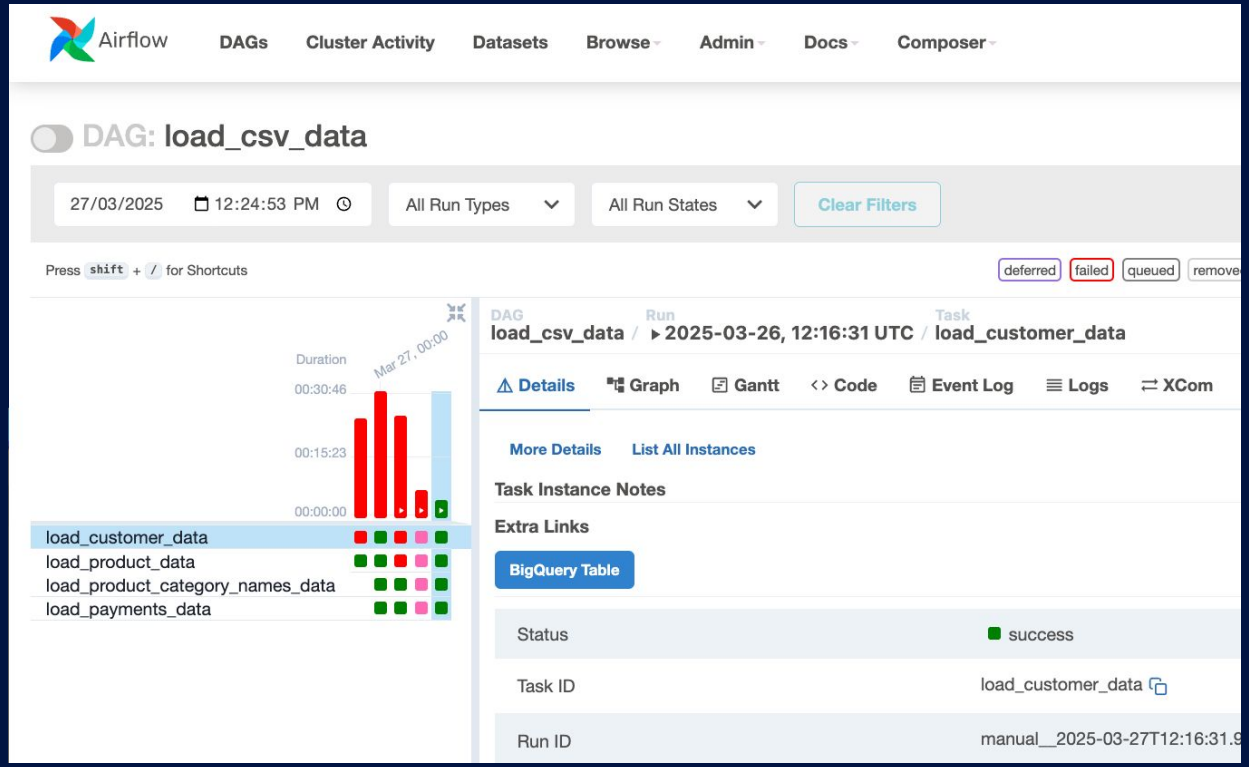

olist Data Engineer Lab

DBT running as DAG

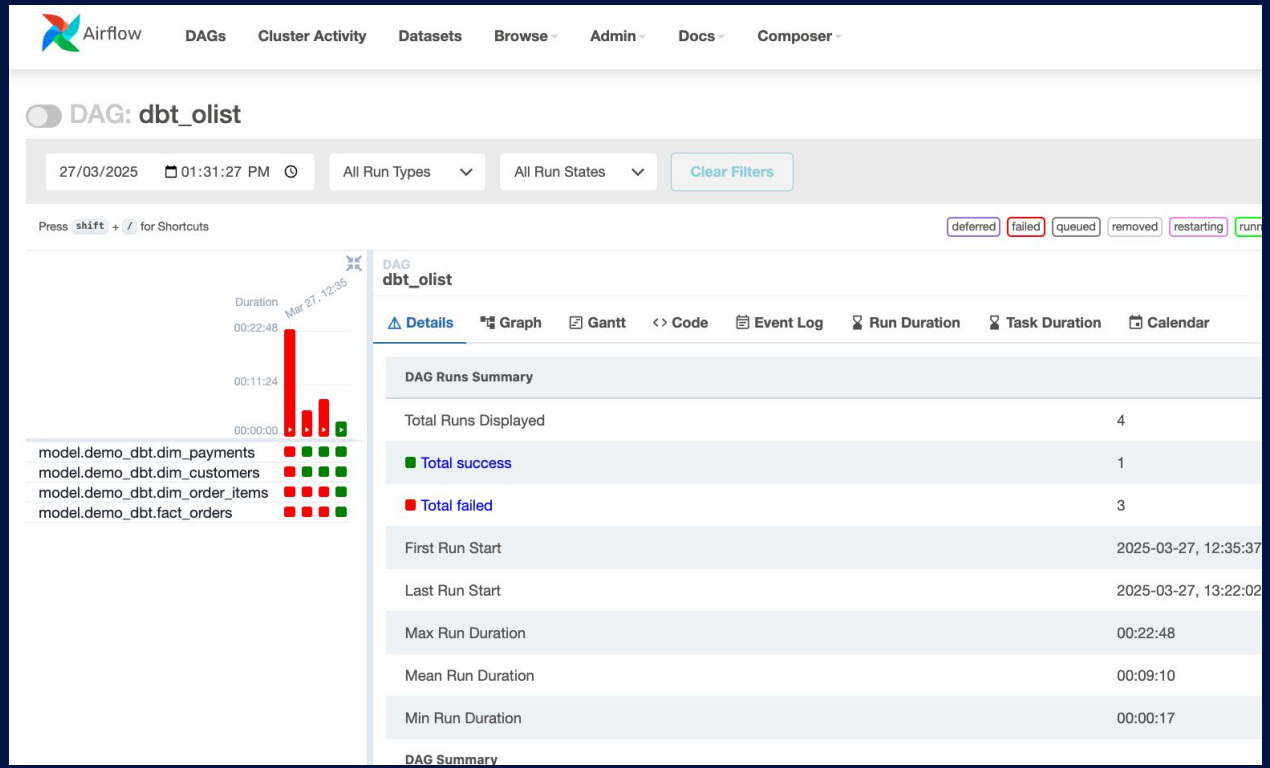

Tables and Views in BigQuery

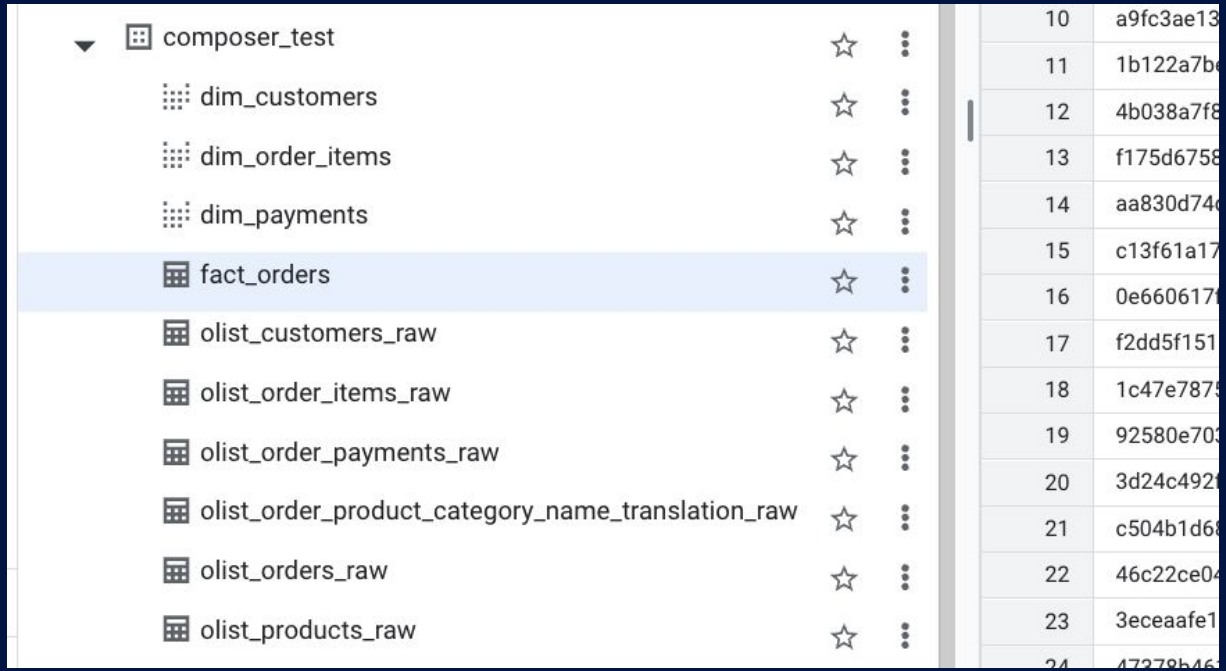

olist Data Engineer Lab

References

https://www.kaggle.com/datasets/olistbr/brazilian-ecommerce/data

https://cloud.google.com/bigquery/docs/samples/bigquery-pandas-gbq-read-gbq-simple

https://stripe.com/en-sg/resources/more/boleto-an-in-depth-guide

https://www.olist.com/

https://docs.getdbt.com/guides/bigquery

https://schemas.getdbt.com/dbt/manifest/v12/index.html#nodes_additionalProperties_anyOf_i4

olist Data Engineer Lab

THANK YOU

+123-456-7890

HELLO@OLIST-DATAELAB.COM

WWW.OLIST-DATAELAB.COM

123 MONSERATT ST., SR, ST 12345